

# CSE 11: Lecture 17

- ✓ Packages
- ✓ Package naming, and CLASSPATH
- ✓ Multithreaded programming
- ✓ Synchronized statements and synchronized methods

Final Exam will be...

**Friday Dec 12 8:00am-10:30am**

Location: here

Closed-book, closed-notes, no calculators. Bring picture ID!

Coverage: Everything...

A practice final is available online (PDF format)

Lecture notes are available online

Final review session will be: Lecture Thursday

# Packages

- ✓ You know that
  - ✗ There are 4 levels of visibility for a method or a variable defined in a class: **public**, **protected**, **private**, and (if none of these is specified) package visibility
  - ✗ Package visibility for a method or a variable means: it is accessible from any method in any other class in the same package
  
- ✓ You also know that
  - ✗ If you don't do anything unusual, all the classes in the same directory are in the same “default” package
  
- ✓ Let's look at Java packages in a little more detail

## Packages, cont'd

- ✓ You can think of a class as a nice way to group some data and functions together
- ✓ Packages are a higher level of software organization: Packages are a means for organizing related classes together
- ✓ In Java, classes in the same package can have special access privileges with respect to one another, and are typically designed to work together, taking advantage of that “package visibility”
- ✓ (This is somewhat similar to so-called “friendly visibility” in C++)
- ✓ If you do not have a package statement declaring package membership at the top of the .java source code file for a class, then that class is part of the unnamed, “default” package, together with other such classes in the same directory
- ✓ Now let's see how to create named packages

## Package names

- ✓ We have seen some package names already, for example

`java.lang`

`java.awt`

`java.awt.event`

- ✓ These are part of the full names of classes found in these packages, for example

`java.lang.String`

`java.awt.Container`

`java.awt.event.ActionEvent`

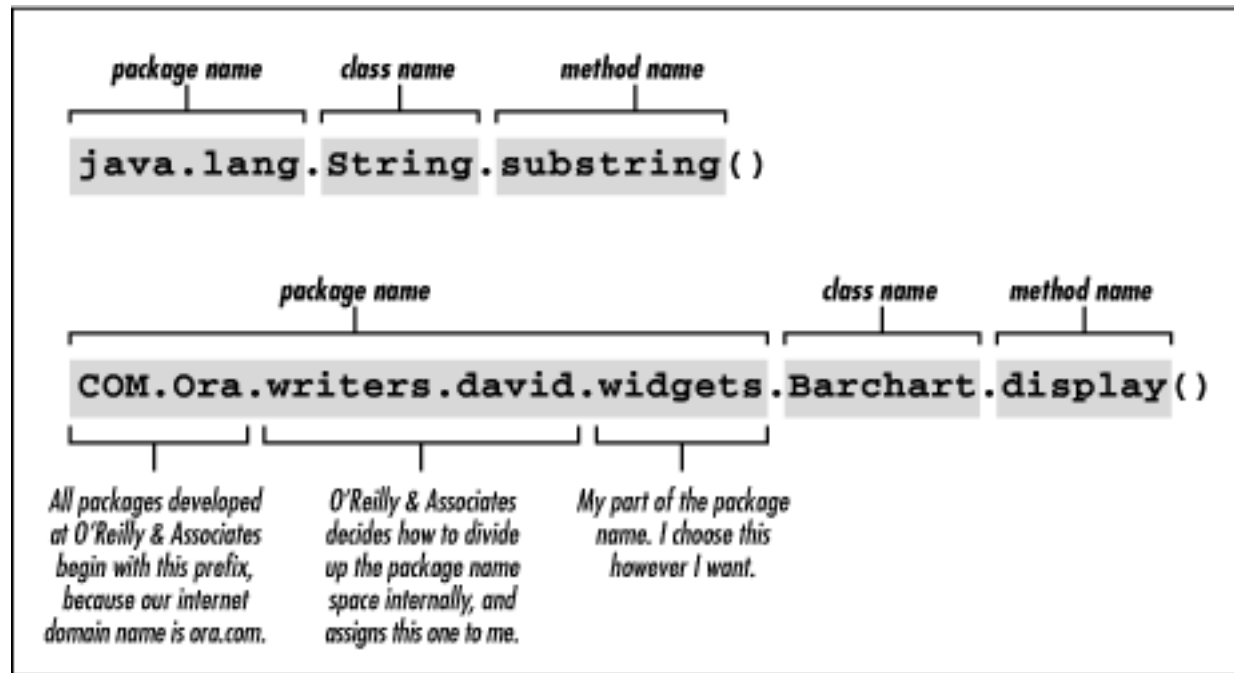
- ✓ Rules for using class names in Java:

- ✗ you do not need to use the full, package-qualified name of classes in the `java.lang` package
- ✗ in a class definition, you do not need to use the package-qualified name of classes in the same package
- ✗ and for other packages, you do not need to use the full, package-qualified name if you import the class using an import statement

## Package names, cont'd

- ✓ Package names are hierarchical: they correspond to the hierarchical directory structure on a machine, and to the Internet domain and host names for the machine
- ✓ On a machine, the package name is relative to paths on the CLASSPATH,
- ✓ CLASSPATH is a list of directories containing .class files; this list by default includes the current working directory, and the parent directory of the standard libraries
  
- ✓ So, for example:
  - ✗ the classes for the `java.lang` package would be in a directory `java/lang` under some directory on the CLASSPATH for your machine
  - ✗ the classes for a package `com.ora.writers.david.widgets` would be in a directory `writers/david/widgets` under some directory on the CLASSPATH on the machine `ora.com`
- ✓ Package naming is part of the beginning of a standard for Uniform Resource Identifiers (URIs) that will provide a persistent identifier for every downloadable object in the world: potentially trillions of them, forming the global infrastructure for all computing resources

## Package names: a picture



## Creating a package, and putting classes in it

- ✓ To create a package, you first create a directory with a certain pathname relative to a directory on the CLASSPATH (e.g. the current working directory):

```
% mkdir foo/bar
```

- ✓ Then put a package statement in each source code file in that directory that declares the contents of the file to be in the corresponding package:

```
package foo.bar;
```

```
public class MyClass {    // etc...
```

- ✓ The package statement must appear as the first statement (i.e., comments can come before it, but nothing else) in the file
- ✓ Now another program that wants to use classes in that package without using their fully qualified name can import them (all of them, with `*`; or individually):

```
import foo.bar.*;
```



# Setting CLASSPATH

- ✓ CLASSPATH is the name of an “environment variable” maintained by the operating system
- ✓ You can check the value of CLASSPATH under Unix (most shells) with either of these commands:

```
printenv CLASSPATH  
echo $CLASSPATH
```

- ✓ If CLASSPATH is undefined, Java will use the current working directory and the Java standard library directories as the CLASSPATH
- ✓ If CLASSPATH is defined, Java will use the Java standard library directories and the directories on the CLASSPATH only
- ✓ You can set the value of CLASSPATH under Unix with the **setenv** command, if you are using the csh shell
- ✓ The CLASSPATH variable should be a list of directory pathnames, separated by **:**'s ... Example: If you want to use the P6 classes without having to copy them to your own directory, you could do this:

```
setenv CLASSPATH .:~/../public/P6
```

# Threads and multithreaded programming

- ✓ In CSE 11, the code you have written has been *single-threaded*: only a single sequence of statements is executed
- ✓ Even when using complicated branching and iteration, your code executes in a single sequence of instructions, i.e., a single thread
- ✓ (Other threads may have been running within the Java Virtual Machine: the garbage collector, the Swing event dispatcher, etc., but your code has been single-threaded)
- ✓ For some purposes, you might want instead to write a *multi-threaded* program:
  - ✗ When working over a network, you might be downloading data from several different machines at once; multithreading would permit doing this in parallel, instead of having to wait until the first download finishes before starting on the next
  - ✗ When doing an animation, it might be convenient to have each character interact with the others as its own thread of execution
  - ✗ etc.
- ✓ Java provides good support for multithreaded programming, which we'll now look at briefly

# Multithreaded programming in Java

- ✓ The key to doing multithreaded programming in Java is the **Thread** class
- ✓ Here are the Thread methods we will consider:

```
public class Thread extends Object implements Runnable {  
  
    public Thread() {...  
  
    public static native void yield() {...  
  
    public void run() {...  
  
    public synchronized native void start() {...
```

- ✓ Now let's look at how to use these methods...

## How to create a thread

- ✓ If you want to do multithreaded programming in Java, the basic way to do it is this:
  - ✗ Define classes that extend Thread, and that override Thread's run() method
  - ✗ Create instances of these classes, and call their start() methods
  - ✗ Calling the start() method of a Thread object does this:
    - It starts a new thread within the Java runtime environment.
    - This new thread executes the statements within the body of the Thread object's run() method.
    - The start() method returns as soon as it starts the thread; the thread keeps running until the run() method returns (or the Thread's stop() method is called, or the thread is interrupted, etc.)
  - ✗ Depending on whether the Java implementation uses preemptive or nonpreemptive thread scheduling, a Thread may have to call yield() in order to let other Threads execute

## Multithreaded programming: an example

```
public class Ones extends Thread {
    public void run () {
        while(true) // just print a sequence of 1's
            System.out.print(1);
    }
}

-----

public class Zeros extends Thread {
    public void run () {
        while(true) // just print a sequence of 0's
            System.out.print(0);
    }
}

-----

public class Test {
    public static void main( String args[] ) {
        Zeros z = new Zeros();
        Ones o = new Ones();
        o.start();
        z.start();
    }
}
```

## An example, cont'd

✓ Q: What happens when that multithreaded program runs?

A: It depends on whether the Java implementation uses preemptive or nonpreemptive thread scheduling

- ✓ Preemptive scheduling: By choosing a maximum length of time a thread can continuously, a preemptive scheduling mechanism guarantees that no single thread uses more than its fair share of the processor. If a thread runs for that amount of time without yielding control to another thread, the scheduler preempts the thread and causes it to stop running so that another thread can run.
- ✓ Nonpreemptive scheduling (also sometimes called “cooperative multitasking”): A nonpreemptive scheduling mechanism cannot preempt threads. A nonpreemptive scheduler relies on the individual threads to yield control of the processor frequently, so that it can provide reasonable performance. A thread explicitly yields control by calling the Thread object’s `yield()` method.

## An example, cont'd

- ✓ So, with preemptive scheduling, the test program might output something like

```
111100011100000111110001111000000011111100000000011111000001110011
11111000001111111000000110010111100000011111100000011111000000001
11111110000000000011111110000111111011111100000001111100100001111
1111001111101111110000000011110000000111111111111 ...
```

- ✓ With nonpreemptive scheduling, the test program will output

```
11111111111111111111111111111111111111111111111111111111111111111111
11111111111111111111111111111111111111111111111111111111111111111111
11111111111111111111111111111111111111111111111111111111111111111111
11111111111111111111111111111111111111111111111111111111111111111111 ...
```

- ✓ Nonpreemptive multithreading doesn't work well for this example!
- ✓ Some platforms may use nonpreemptive scheduling for Java threads. To get reasonable threadswitching on any platform, we make the threads call `yield()` periodically...

## Multithreaded programming: a revised example

```
public class Ones extends Thread {
    public void run () {
        long count = 0;
        while(true) { // just print a sequence of 1's
            System.out.print(1);
            if(++count % 10 == 0) yield(); // yield every 10 steps
        }
    }
}

-----

public class Zeros extends Thread {
    public void run () {
        long count = 0;
        while(true) { // just print a sequence of 0's
            System.out.print(0);
            if(++count % 10 == 0) yield(); // yield every 10 steps
        }
    }
}
```



## A revised example, cont'd

✓ Now, even with nonpreemptive scheduling, the test program will output

```
11111111110000000000111111111100000000001111111111000000000011111111
11100000000000111111111100000000001111111111000000000011111111110000
0000001111111111000000000011111111110000000000111111111100000000001
111111111000000000011111111100000000001111111110000000000111111111
1100000000001111111111000000000011111111110000000000 ...
```

## Synchronization and locks

- ✓ Each Java virtual machine can support many threads of execution at once
- ✓ These threads independently execute code that operates on objects residing in a shared main memory
  - ✗ (This is the main difference between multithreading and multiprocessing: Processes do not share memory, but multiple threads within the same process do.)
- ✓ How to make sure that these independent actions will produce a nice result? Answer: *synchronization*
- ✓ To synchronize threads, every Java object has a *lock* associated with it
- ✓ Only one thread can acquire the lock for an object at any one time; no other thread can acquire the lock until the lock is released by the thread that has the lock
- ✓ *Synchronized statements and methods* are the way to acquire and release an object's lock

## Synchronized statements

- ✓ A synchronized statement has the form

```
synchronized (<expr>) {  
    <body>  
}
```

- ✓ The synchronized statement has this semantics:

- ✗ First, the expression <expr> is evaluated. It must evaluate to a non-null object reference
- ✗ Next, the lock on that object is acquired. If the lock is already owned by another thread, the current thread stops executing and waits for it to be released
- ✗ Then, the statements in the body of the synchronized statement are executed
- ✗ Finally, the lock is released.

- ✓ Note: nothing prevents another thread from modifying the “locked” object, if that other thread is not synchronized on it!
- ✓ Synchronization is essential for ensuring deterministic behavior in multithreaded programming, but the programmer has to understand when and how to use it...

## Synchronized methods

- ✓ An instance method can be declared synchronized
- ✓ This is exactly the same as declaring that the body of the method is synchronized on the “this” object
- ✓ For example:

```
public synchronized void setTime(int hour, int min, int sec) {  
    h = hour; m = min; s = sec;  
}
```

is the same as

```
public void setTime(int hour, int min, int sec) {  
    synchronized(this) { h = hour; m = min; s = sec; }}
```

- ✓ Without synchronization here, two threads simultaneously executing the setTime method of an object might put the object in an inconsistent state
- ✓ In general, methods that might not work correctly if executed by two threads simultaneously should be declared synchronized, or they should be documented as not being thread-safe

## Next time

- ✓ Compile-time and run-time type checking
- ✓ Type checking and dynamic method binding
- ✓ Final review of CSE 11 topics